

How you ask changes how it reasons, not what it predicts

Abdul Ahad



UNIVERSITÄT
DES
SAARLANDES

Question

English marks confidence in the verb. *Predict* commits, while *guess* is more open-ended. So if a model is pragmatically sensitive, then switching verbs should change the answer.

Three questions

Q1 Does the epistemic verb change accuracy? (*no*)

Q2 Is it the verb, or the question? (*neither*)

Q3 Then what does change? (*how it reasons*)

Paradigm

Rock Paper Scissors, 100 rounds, scripted opponent. Labels **J**, **F**, **M**, so no memorised strategy transfers. Qwen3.5-9B, vLLM.

You are playing 100 rounds against the same opponent.
Each round you both pick one of: J, F, M.
F beats J, M beats F, J beats M. Same pick is a tie.
...

Round 1
Opponent move: F
Your move: J
...

The opponent's move depends on what you played in the previous round.
You are currently playing round 4.

What do you predict the other player will play in this round?


Opponent move: _

Two lines move; every other byte is identical.

- **The ask:** 10 verbs, high (*predict*) to low (*guess*), $\times 4$ forms, question to imperative.
- **The task line:** absent, or a pointer that is true for mirror and tft, false for cycle and random.

Opponents: cycle, mirror, tft, counter and random (*the falsifier: nothing to predict*).

The verb does not change accuracy



- Ten verbs ranked in 126 situations, counting firsts minus lasts.
- Nothing escapes the chance band, so there is no verb ordering that holds across the board.
- **One Exception:** on mirror with no hint, high confidence verbs did beat low ones, 45.5% against 30.2%, and it held in every phrasing.
- But the same opponent in another setup reverses it, and the 126 cell test finds nothing. So the effect is real where it appears and does not travel.

What the verb does change: the reasoning


Accuracy is flat, but the register replicates across opponents.

High verbs reason in the language of rules; mid and low verbs in the language of likelihood,

“the question asks to infer, suggesting a specific answer”

“estimate implies a probabilistic guess”


mirror, tft and counter need cross stream reasoning: your column against theirs at a one round lag.



lag 1 pairs written per response, mirror L0

- *infer* writes twice as many correctly lagged pairs as any other verb.
- Doing it decides the answer: **84.7%** right when the model does it, **32.9%** when it does not, yet your own column is almost never the first hypothesis.

One true sentence beats any verb



- A true hint buys 33 to 47 points, a false one costs 52. On *random*, false but with nothing to predict, it does nothing.
- *cycle* is solved at **96.2%** without the hint, **44.6%** with it: the model follows a false instruction against a history that contradicts it.

What this means

- **The pragmatic hypothesis fails on accuracy**, and so does the interrogative frame. But the wording is not inert: it changes the search strategy.
- Telling the model where to look is worth far more than any word choice, and telling it wrongly is worse than saying nothing.
- The model can find the right rule and still fail to hold on to it.

What is running now

The wider RPS experiments all used *predict*, which the evidence above makes a non neutral choice. A replication with *infer* is running.